
Claude Code: What's Coming Next

8 unreleased and rolling-out features that will change how you build.
What each one does, why it matters, and exactly what to do before they go live.

APRIL 2026 / BISHOP AI INTELLIGENCE BRIEF

FEATURE 01

KAIROS

An always-on background daemon that watches your repos 24/7 and acts before you ask.

WHY THIS MATTERS


Claude stops being a tool you invoke and becomes a persistent teammate. It subscribes to GitHub webhooks, triages issues, monitors PRs, and prepares context before you sit down. Found 150+ references in leaked source code with a May 2026 target.

PREPARE NOW

Ensure your repos have webhook access configured. Document every recurring manual task you do weekly: PR reviews, dependency checks, log monitoring, status updates. These become KAIROS handoffs on day one.

ACTION PATH

1. Audit your weekly manual tasks → 2. Set up GitHub webhook access → 3. Write CLAUDE.md instructions for each task → 4. Enable KAIROS on launch day



FEATURE 02

AutoDream

Memory consolidation that runs while you sleep. Wake up to a sharper, cleaner agent.

WHY THIS MATTERS


A forked subagent activates during idle periods, merging observations, removing contradictions, and converting vague insights into verified facts. It gets read-only bash access to confirm things against real code. Your next session starts with pristine context instead of cluttered reasoning.

PREPARE NOW

Start using Claude Code's memory system heavily. The more raw material AutoDream has, the better it consolidates. Write detailed CLAUDE.md files. Leave sessions idle instead of killing them.

ACTION PATH

1. Populate your memory directory with project context → 2. Write thorough CLAUDE.md → 3. Let sessions idle instead of force-closing → 4. Review consolidated memories on launch



FEATURE 03

ULTRAPLAN

Offload complex architectural decisions to a remote cloud session with 30 minutes of dedicated think time.

WHY THIS MATTERS


Migration plans, system redesigns, and complex refactoring strategies get dedicated Opus 4.6 deep-thinking time in a cloud container. Includes a browser-based approval workflow so you don't accidentally trigger massive compute bills. Your local session stays unblocked.

PREPARE NOW

Document your hardest architectural problems in structured markdown. The clearer your input, the better the 30-minute deep-think produces. Budget for a new cost tier since extended Opus thinking will be premium.

ACTION PATH

1. Write up your top 3 architectural decisions → 2. Structure them as problem + constraints + options → 3. Set a compute budget → 4. Submit via ULTRAPLAN on launch



FEATURE 04

Coordinator Mode

A structured orchestrator managing parallel worker agents with defined roles and formal handoffs.

WHY THIS MATTERS


Agent Teams is experimental and free-form. Coordinator Mode adds structure: named roles, defined handoff protocols, progress tracking, and conflict resolution. Think of it as the difference between a group chat and a managed sprint. Multi-agent workflows become reliable enough for production.

PREPARE NOW

If you use Agent Teams today, document which patterns work. Identify tasks that need parallel execution with coordination: full-stack features, multi-repo changes, large refactors. Design your ideal agent team structure now.

ACTION PATH

1. Map your multi-agent workflows →
2. Define roles (frontend, backend, tests, review) →
3. Write handoff criteria for each role →
4. Migrate from Agent Teams on launch



FEATURE 05

Auto Mode

An AI classifier silently approves tool permissions. Zero confirmation prompts. Full autonomy.

WHY THIS MATTERS


Today, permission prompts interrupt flow. Sandbox mode reduced them by 84%, but they still appear for network access, file writes outside CWD, and destructive commands. Auto Mode uses a classifier to evaluate each request and approve or deny without human intervention. Claude becomes fully autonomous.

PREPARE NOW

Get comfortable with sandbox mode first since Auto Mode builds on it. Audit your CLAUDE.md for safety boundaries. Only plan to use Auto Mode in isolated environments until you trust the classifier's judgment.

ACTION PATH

1. Enable and test sandbox mode today → 2. Tighten your CLAUDE.md safety rules → 3. Set up isolated dev environments → 4. Enable Auto Mode in sandboxed repos first



FEATURE 06

Voice Mode

Hold spacebar to talk. Mix typed and spoken input in a single prompt. 20 languages.

WHY THIS MATTERS


Describing complex architectural changes is 3x faster by voice than typing. Voice Mode started rolling out to ~5% of users in March 2026 with broader availability ongoing. You can mix voice and typed input in the same prompt, making it natural to narrate intent while pointing at specific code.

PREPARE NOW

Make sure you're on Pro or Max plan. Practice describing code changes verbally. Think about which workflows are bottlenecked by typing speed: code reviews, architectural explanations, rapid prototyping instructions.

ACTION PATH

1. Update to latest Claude Code version → 2. Check if voice mode is available in your account → 3. Start with code review narration → 4. Expand to architectural design sessions



FEATURE 07

Agent Teams

2 to 16 Claude sessions working your codebase in parallel. One leads, the rest execute.

WHY THIS MATTERS

Currently experimental behind a feature flag. Teammates share findings directly via mailbox, not through a central agent. A full-stack feature that takes sequential hours can run frontend, backend, and tests simultaneously. Expect 3-4x token usage but massive time compression.

PREPARE NOW

Set `CLAUDE_CODE_EXPERIMENTAL_AGENT_TEAMS=1` and test today. Structure your repos for parallel work: clear module boundaries, independent test suites, well-defined interfaces between components.

ACTION PATH

1. Enable the feature flag now → 2. Identify a multi-file feature to test with → 3. Define team roles in your prompt → 4. Scale to larger teams as patterns stabilize



FEATURE 08

Claude Mythos

Codename "Capybara." The next model family, reportedly dramatically better at programming.

WHY THIS MATTERS


Internal testing shows Mythos outperforms Opus 4.6 significantly on programming benchmarks. Comes with a rumored "Dev Team" mode that auto-generates sub-agents for collaborative development. Every skill, prompt, and workflow you've built will get an instant quality boost the day you switch models.

PREPARE NOW

Invest in your CLAUDE.md files, skill library, and memory system now. These are model-agnostic infrastructure. When Mythos drops, better prompts + better model = compounding returns. Don't wait for the upgrade to build your tooling.

ACTION PATH

1. Perfect your CLAUDE.md and skills today → 2. Build reusable prompt templates → 3. Set up evaluation benchmarks for your workflows → 4. Swap model ID on launch day and measure the lift



Your Action Plan: Do This Now

- 1 Update Claude Code to v2.1.91+**
Everything downstream depends on being current. Run `npm update -g @anthropic-ai/claude-code` today.
- 2 Document your recurring manual tasks**
PR reviews, dependency checks, log monitoring, status reports. These become KAIROS handoffs on launch day.
- 3 Enable Agent Teams flag now**
Set `CLAUDE_CODE_EXPERIMENTAL_AGENT_TEAMS=1` and start testing parallel workflows before Coordinator Mode ships.
- 4 Invest in your memory + CLAUDE.md**
AutoDream and Mythos both amplify what you've already built. Better inputs = compounding returns.
- 5 Enable sandbox mode**
Foundation for Auto Mode. Test it in your dev environment so you're ready when full autonomy ships.
- 6 Structure architectural problems**
Write them as problem + constraints + options. Ready for ULTRAPLAN's 30-minute deep-think sessions.
- 7 Set up evaluation benchmarks**
Measure your current workflow speed and quality. When Mythos drops, you'll know exactly how much it improved.
- 8 Watch May 2026**
KAIROS, AutoDream, ULTRAPLAN, and Coordinator Mode are all targeting launch. Be ready, not scrambling.